Instance A (gradient descent) validation learning curves — Run 1, real checkpoints

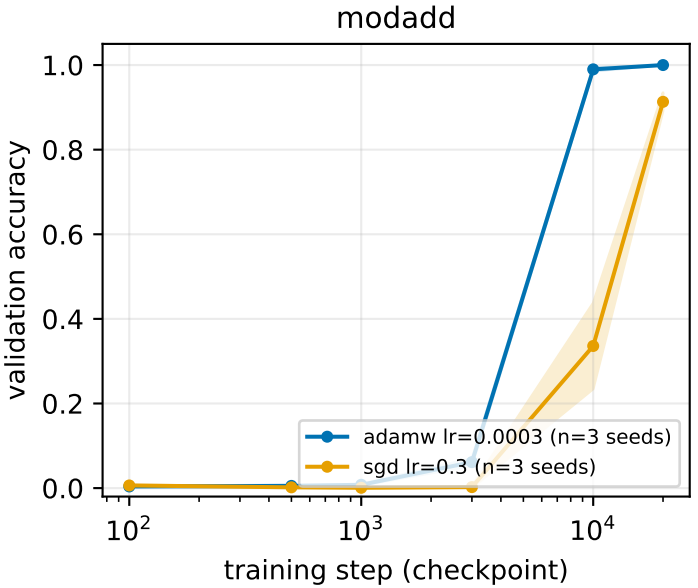

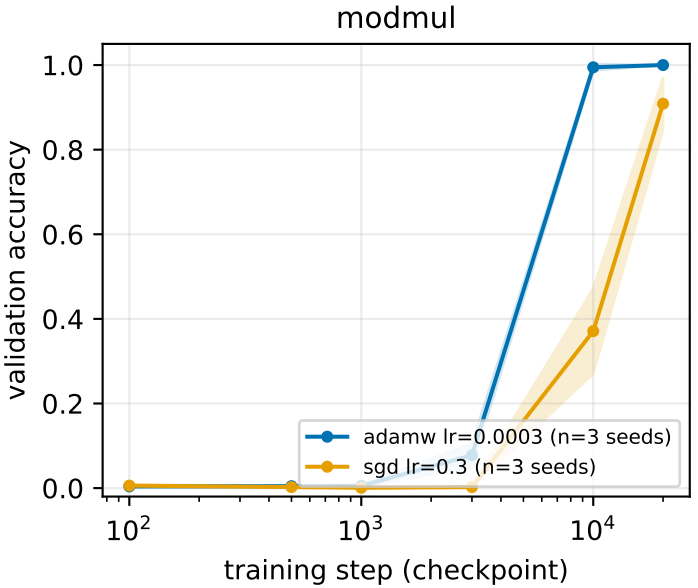

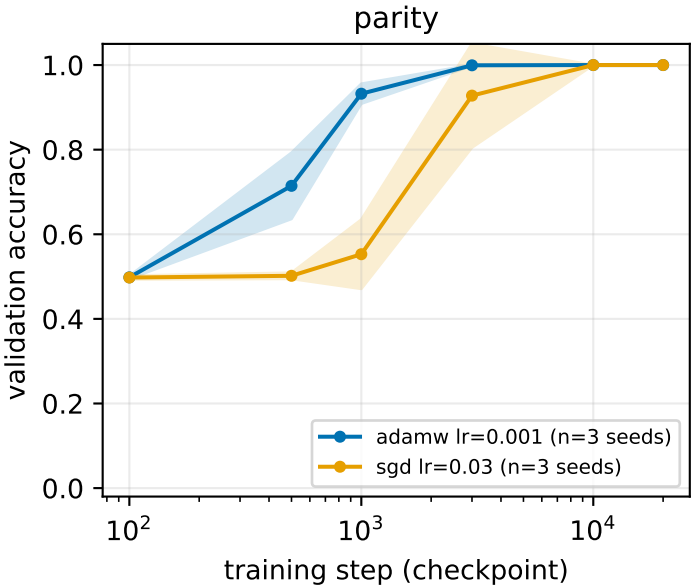

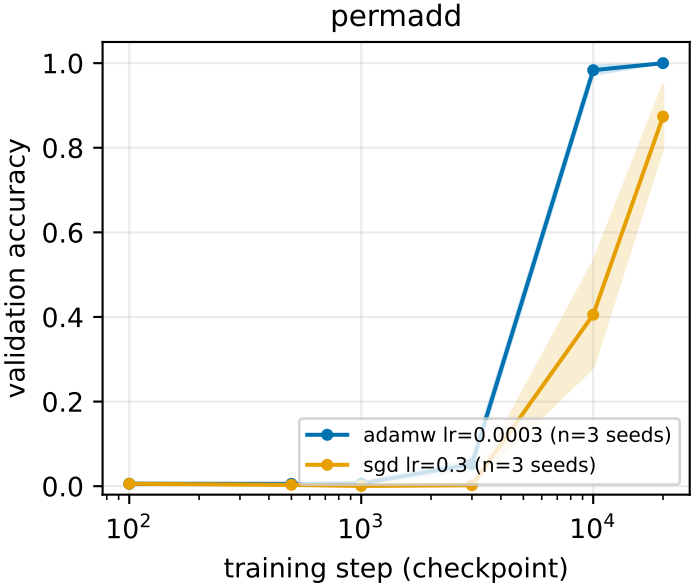

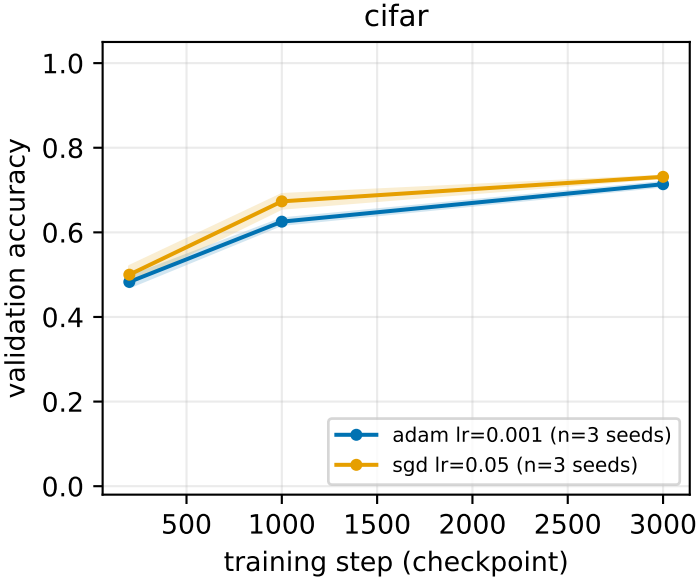

Source: v3_run1_outputs/A_curves.csv
Algo split train/val/test = 0.3/0.1/0.6; seeds 0,1,2; checkpoints as logged.
LR shown = argmax mean final val_acc among the Run-1 grid (not interpolated).
Run 2 / local-pilot A_curves.csv empty (A disabled)—no fabricated curves.